

feature-selection

October 16, 2023

```
[1]: # Loading Libraries
import pandas as pd
import numpy as np
from matplotlib import pyplot as plt
import seaborn as sns
```

```
[2]: # Loading Dataframe
df = pd.read_csv('mobile_data.csv')
```

```
[3]: df.head()
```

```
[3]:  battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  m_dep  \
0           772     0         1.1         1  12      0          39    0.8
1          1709     1         2.1         0   1      0          13    1.0
2          1949     0         2.6         1   4      0          47    0.3
3           615     1         2.5         0   0      0          10    0.8
4          1821     1         1.2         0  13      1          44    0.6
```

```
    mobile_wt  n_cores  ...  px_height  px_width  ram  sc_h  sc_w  talk_time  \
0          81         7  ...    1314    1854  2819   17   15          3
1         156         2  ...     974    1385  3283   17    1         15
2         199         4  ...     407     822  1433   11    5         20
3         131         6  ...    1216    1786  2769   16    8         11
4         141         2  ...    1208    1212  1411    8    2         15
```

```
    three_g  touch_screen  wifi  price_range
0         1             1     0             3
1         1             0     0             3
2         0             0     1             1
3         1             0     0             2
4         1             1     0             1
```

[5 rows x 21 columns]

```
[4]: # Dataframe summary
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 2000 entries, 0 to 1999

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
0	battery_power	2000 non-null	int64
1	blue	2000 non-null	int64
2	clock_speed	2000 non-null	float64
3	dual_sim	2000 non-null	int64
4	fc	2000 non-null	int64
5	four_g	2000 non-null	int64
6	int_memory	2000 non-null	int64
7	m_dep	2000 non-null	float64
8	mobile_wt	2000 non-null	int64
9	n_cores	2000 non-null	int64
10	pc	2000 non-null	int64
11	px_height	2000 non-null	int64
12	px_width	2000 non-null	int64
13	ram	2000 non-null	int64
14	sc_h	2000 non-null	int64
15	sc_w	2000 non-null	int64
16	talk_time	2000 non-null	int64
17	three_g	2000 non-null	int64
18	touch_screen	2000 non-null	int64
19	wifi	2000 non-null	int64
20	price_range	2000 non-null	int64

dtypes: float64(2), int64(19)

memory usage: 328.2 KB

```
[5]: # total null value
df.isnull().sum()
```

```
[5]: battery_power    0
blue                0
clock_speed         0
dual_sim            0
fc                  0
four_g              0
int_memory          0
m_dep               0
mobile_wt           0
n_cores             0
pc                  0
px_height           0
px_width            0
ram                 0
sc_h                0
sc_w                0
```

```

talk_time      0
three_g        0
touch_screen   0
wifi           0
price_range    0
dtype: int64

```

```
[6]: x = df.drop('price_range',axis=1)
     y = df.price_range
```

```
[7]: x.head()
```

```
[7]:  battery_power  blue  clock_speed  dual_sim  fc  four_g  int_memory  m_dep  \
0           772     0         1.1         1  12      0         39    0.8
1          1709     1         2.1         0   1      0         13    1.0
2          1949     0         2.6         1   4      0         47    0.3
3           615     1         2.5         0   0      0         10    0.8
4          1821     1         1.2         0  13      1         44    0.6

    mobile_wt  n_cores  pc  px_height  px_width  ram  sc_h  sc_w  talk_time  \
0          81        7  14      1314      1854  2819   17   15          3
1         156        2   2       974      1385  3283   17    1         15
2         199        4   7       407       822  1433   11    5         20
3         131        6   9      1216      1786  2769   16    8         11
4         141        2  14      1208      1212  1411    8    2         15

    three_g  touch_screen  wifi
0         1             1     0
1         1             0     0
2         0             0     1
3         1             0     0
4         1             1     0

```

```
[8]: y.head()
```

```
[8]: 0    3
     1    3
     2    1
     3    2
     4    1
     Name: price_range, dtype: int64

```

```
[9]: y.value_counts()
```

```
[9]: 3    501
     1    501
     0    500

```

2 498
Name: price_range, dtype: int64

```
[10]: # Splitting dataframe
from sklearn.model_selection import train_test_split
xtrain,xtest,ytrain,ytest = train_test_split(x,y,test_size=.3,random_state=43)
```

```
[11]: from sklearn.tree import DecisionTreeClassifier
```

0.0.1 Checking Fittings

```
[12]: acc_training = []
acc_testing = []
depth = [ n for n in range(1,20)]
for i in depth:
    clf = DecisionTreeClassifier(max_depth = i)
    clf.fit(xtrain,ytrain)
    score_train = clf.score(xtrain,ytrain)
    acc_training.append(score_train)
    score_test = clf.score(xtest,ytest)
    acc_testing.append(score_test)


    print('Depth = %d , Train = %f , Test = %f ' % (i,score_train,score_test) )
```

```
Depth = 1 , Train = 0.492857 , Test = 0.518333
Depth = 2 , Train = 0.760000 , Test = 0.763333
Depth = 3 , Train = 0.766429 , Test = 0.766667
Depth = 4 , Train = 0.841429 , Test = 0.793333
Depth = 5 , Train = 0.888571 , Test = 0.820000
Depth = 6 , Train = 0.917143 , Test = 0.828333
Depth = 7 , Train = 0.947143 , Test = 0.830000
Depth = 8 , Train = 0.974286 , Test = 0.818333
Depth = 9 , Train = 0.987857 , Test = 0.831667
Depth = 10 , Train = 0.995714 , Test = 0.823333
Depth = 11 , Train = 0.999286 , Test = 0.820000
Depth = 12 , Train = 1.000000 , Test = 0.821667
Depth = 13 , Train = 1.000000 , Test = 0.831667
Depth = 14 , Train = 1.000000 , Test = 0.816667
Depth = 15 , Train = 1.000000 , Test = 0.825000
Depth = 16 , Train = 1.000000 , Test = 0.823333
Depth = 17 , Train = 1.000000 , Test = 0.825000
Depth = 18 , Train = 1.000000 , Test = 0.811667
Depth = 19 , Train = 1.000000 , Test = 0.825000
```

```
[13]: plt.plot(depth,acc_training,"-o",label='Training')
plt.plot(depth,acc_testing,"-o",label='Testing')
plt.xlabel('Depth')
```

```
plt.ylabel('Accuracy')
plt.legend()
```

[13]: <matplotlib.legend.Legend at 0x284670246c8>



From the plot, we can see all are going smoothly and not having much difference between them, so they are good fitting

0.1 Feature Selection

0.1.1 Extra Trees Classifier

```
[14]: from sklearn.ensemble import ExtraTreesClassifier
      etc = ExtraTreesClassifier()
```

```
[15]: etc.fit(x,y)
```

[15]: ExtraTreesClassifier()

```
[16]: # Gini Feature importance
      feature_importance = etc.feature_importances_
      feature_importance #info gain
```

```
[16]: array([0.06180858, 0.01968337, 0.03293631, 0.01981819, 0.03189973,
            0.0174042 , 0.03388182, 0.03347493, 0.03604542, 0.03242053,
            0.03242006, 0.04742122, 0.04848003, 0.39814073, 0.0338445 ,
            0.034256 , 0.03390626, 0.01479553, 0.01786608, 0.01949652])
```

```
[17]: imp = pd.DataFrame(feature_importance, columns=['Gain_Score'])
cols = pd.DataFrame(x.columns, columns=['Feature_Names'])
gains_x = pd.concat([cols,imp],axis=1)
gains_x
```

```
[17]:
```

	Feature_Names	Gain_Score
0	battery_power	0.061809
1	blue	0.019683
2	clock_speed	0.032936
3	dual_sim	0.019818
4	fc	0.031900
5	four_g	0.017404
6	int_memory	0.033882
7	m_dep	0.033475
8	mobile_wt	0.036045
9	n_cores	0.032421
10	pc	0.032420
11	px_height	0.047421
12	px_width	0.048480
13	ram	0.398141
14	sc_h	0.033845
15	sc_w	0.034256
16	talk_time	0.033906
17	three_g	0.014796
18	touch_screen	0.017866
19	wifi	0.019497

```
[18]: gains_x.nlargest(10, 'Gain_Score')
```

```
[18]:
```

	Feature_Names	Gain_Score
13	ram	0.398141
0	battery_power	0.061809
12	px_width	0.048480
11	px_height	0.047421
8	mobile_wt	0.036045
15	sc_w	0.034256
16	talk_time	0.033906
6	int_memory	0.033882
14	sc_h	0.033845
7	m_dep	0.033475

```
[19]: gains_x.nsmallest(10, 'Gain_Score')
```

```
[19]:
```

	Feature_Names	Gain_Score
17	three_g	0.014796
5	four_g	0.017404
18	touch_screen	0.017866


```

19      wifi      0.019497
1       blue      0.019683
3      dual_sim   0.019818
4       fc        0.031900
10      pc        0.032420
9      n_cores    0.032421
2  clock_speed    0.032936

```

```
[21]: gains_x.plot(kind='bar',figsize=(15,10))
```

```
[21]: <matplotlib.axes._subplots.AxesSubplot at 0x28468327508>
```



```
[98]: features = pd.Series(etc.feature_importances_, index = x.columns)
      features
```

```

[98]: battery_power    0.061809
      blue             0.019683
      clock_speed      0.032936
      dual_sim         0.019818
      fc               0.031900
      four_g           0.017404
      int_memory        0.033882
      m_dep             0.033475
      mobile_wt         0.036045

```

```

n_cores      0.032421
pc            0.032420
px_height     0.047421
px_width      0.048480
ram           0.398141
sc_h          0.033845
sc_w          0.034256
talk_time     0.033906
three_g       0.014796
touch_screen  0.017866
wifi          0.019497
dtype: float64

```

```


[23]: features.nlargest(15).plot(kind='barh',figsize=(15,10))
      #plt.savefig('score1.png')

```

```

[23]: <matplotlib.axes._subplots.AxesSubplot at 0x28468720788>

```



```

[24]: # Feature correlations
      corr = x.corr()
      corr

```

```

[24]:
      battery_power      blue  clock_speed  dual_sim      fc \
battery_power      1.000000  0.011792    0.012088 -0.041499  0.031728
blue                0.011792  1.000000    0.022208  0.033198  0.004421

```


clock_speed	0.012088	0.022208	1.000000	-0.000013	-0.002228
dual_sim	-0.041499	0.033198	-0.000013	1.000000	-0.027660
fc	0.031728	0.004421	-0.002228	-0.027660	1.000000
four_g	0.016252	0.010430	-0.040497	0.002204	-0.016691
int_memory	-0.003256	0.040035	0.008211	-0.016619	-0.028369
m_dep	0.035329	0.004904	-0.013967	-0.024364	-0.001180
mobile_wt	0.003405	-0.008509	0.013040	-0.008833	0.023103
n_cores	-0.029628	0.036176	-0.005288	-0.025111	-0.014089
pc	0.029055	-0.009759	-0.006902	-0.016035	0.644736
px_height	0.015460	-0.008042	-0.012916	-0.023092	-0.010013
px_width	-0.006547	-0.043692	-0.006062	0.011648	-0.005447
ram	-0.001211	0.025296	0.003320	0.041313	0.015840
sc_h	-0.029862	-0.002829	-0.028834	-0.012072	-0.009773
sc_w	-0.020972	0.002223	-0.008453	-0.014825	-0.011747
talk_time	0.050825	0.015683	-0.014586	-0.037682	-0.005679
three_g	0.011937	-0.032583	-0.044436	-0.014008	0.002206
touch_screen	-0.011438	0.009071	0.019796	-0.018137	-0.013414
wifi	-0.008686	-0.019863	-0.025748	0.022740	0.018552

	four_g	int_memory	m_dep	mobile_wt	n_cores	pc	\
battery_power	0.016252	-0.003256	0.035329	0.003405	-0.029628	0.029055	
blue	0.010430	0.040035	0.004904	-0.008509	0.036176	-0.009759	
clock_speed	-0.040497	0.008211	-0.013967	0.013040	-0.005288	-0.006902	
dual_sim	0.002204	-0.016619	-0.024364	-0.008833	-0.025111	-0.016035	
fc	-0.016691	-0.028369	-0.001180	0.023103	-0.014089	0.644736	
four_g	1.000000	0.006831	-0.004381	-0.015238	-0.030379	-0.005887	
int_memory	0.006831	1.000000	0.006267	-0.033450	-0.028415	-0.033384	
m_dep	-0.004381	0.006267	1.000000	0.021180	-0.002929	0.026722	
mobile_wt	-0.015238	-0.033450	0.021180	1.000000	-0.018178	0.018626	
n_cores	-0.030379	-0.028415	-0.002929	-0.018178	1.000000	-0.002329	
pc	-0.005887	-0.033384	0.026722	0.018626	-0.002329	1.000000	
px_height	-0.021476	0.008719	0.025173	0.001784	-0.007519	-0.018958	
px_width	0.005709	-0.010383	0.022626	0.001767	0.024629	0.003140	
ram	0.007835	0.032136	-0.010876	-0.003159	0.004643	0.030231	
sc_h	0.025434	0.037661	-0.024976	-0.033877	0.000039	0.005393	
sc_w	0.037128	0.013886	-0.017654	-0.021301	0.026433	-0.023592	
talk_time	-0.045850	-0.001618	0.017614	0.004002	0.013272	0.016714	
three_g	0.583661	-0.010301	-0.014169	0.002776	-0.015518	-0.001586	
touch_screen	0.014719	-0.028666	-0.003156	-0.014787	0.023113	-0.007426	
wifi	-0.016604	0.007938	-0.026069	-0.000497	-0.009535	0.004197	

	px_height	px_width	ram	sc_h	sc_w	talk_time	\
battery_power	0.015460	-0.006547	-0.001211	-0.029862	-0.020972	0.050825	
blue	-0.008042	-0.043692	0.025296	-0.002829	0.002223	0.015683	
clock_speed	-0.012916	-0.006062	0.003320	-0.028834	-0.008453	-0.014586	
dual_sim	-0.023092	0.011648	0.041313	-0.012072	-0.014825	-0.037682	
fc	-0.010013	-0.005447	0.015840	-0.009773	-0.011747	-0.005679	


four_g	-0.021476	0.005709	0.007835	0.025434	0.037128	-0.045850
int_memory	0.008719	-0.010383	0.032136	0.037661	0.013886	-0.001618
m_dep	0.025173	0.022626	-0.010876	-0.024976	-0.017654	0.017614
mobile_wt	0.001784	0.001767	-0.003159	-0.033877	-0.021301	0.004002
n_cores	-0.007519	0.024629	0.004643	0.000039	0.026433	0.013272
pc	-0.018958	0.003140	0.030231	0.005393	-0.023592	0.016714
px_height	1.000000	0.509613	-0.020823	0.059052	0.043935	-0.009059
px_width	0.509613	1.000000	0.003001	0.020039	0.035830	0.007063
ram	-0.020823	0.003001	1.000000	0.015294	0.034836	0.011741
sc_h	0.059052	0.020039	0.015294	1.000000	0.507638	-0.016224
sc_w	0.043935	0.035830	0.034836	0.507638	1.000000	-0.023888
talk_time	-0.009059	0.007063	0.011741	-0.016224	-0.023888	1.000000
three_g	-0.033655	-0.001594	0.016053	0.011672	0.032786	-0.041589
touch_screen	0.020292	-0.004742	-0.029985	-0.020450	0.013447	0.020197
wifi	0.054076	0.033064	0.022397	0.026037	0.033550	-0.031255

	three_g	touch_screen	wifi
battery_power	0.011937	-0.011438	-0.008686
blue	-0.032583	0.009071	-0.019863
clock_speed	-0.044436	0.019796	-0.025748
dual_sim	-0.014008	-0.018137	0.022740
fc	0.002206	-0.013414	0.018552
four_g	0.583661	0.014719	-0.016604
int_memory	-0.010301	-0.028666	0.007938
m_dep	-0.014169	-0.003156	-0.026069
mobile_wt	0.002776	-0.014787	-0.000497
n_cores	-0.015518	0.023113	-0.009535
pc	-0.001586	-0.007426	0.004197
px_height	-0.033655	0.020292	0.054076
px_width	-0.001594	-0.004742	0.033064
ram	0.016053	-0.029985	0.022397
sc_h	0.011672	-0.020450	0.026037
sc_w	0.032786	0.013447	0.033550
talk_time	-0.041589	0.020197	-0.031255
three_g	1.000000	0.012131	0.004316
touch_screen	0.012131	1.000000	0.012904
wifi	0.004316	0.012904	1.000000

```
[25]: # feature_names = x.columns
# feature_names = x.corr().index

plt.figure(figsize=(20,20))
sns.heatmap(corr,annot=True)
```

```
[25]: <matplotlib.axes._subplots.AxesSubplot at 0x28468469488>
```



0.1.2 Select K-Best

```
[26]: # chi2 for non-negative feature
      # f_regression for regression
      # f_classif for classification set
      from sklearn.feature_selection import SelectKBest, f_regression, f_classif, chi2
      skb = SelectKBest(score_func = f_classif)
```

```
[27]: skb.fit(x,y)
```

```
[27]: SelectKBest()
```

```
[29]: # Scores of features
feature_importance = skb.scores_
feature_importance
```

```
[29]: array([3.11908732e+01, 4.59678478e-01, 6.13019154e-01, 4.89492949e-01,
        8.30396133e-01, 1.16710790e+00, 2.96575777e+00, 1.64410175e+00,
        3.62521255e+00, 2.58284723e+00, 8.74295515e-01, 1.95473713e+01,
        2.25200970e+01, 3.52623236e+03, 2.20350431e+00, 1.59731005e+00,
        1.66657575e+00, 4.39621377e-01, 1.45891088e+00, 2.61677203e-01])
```

```
[30]: imp = pd.DataFrame(feature_importance,columns=['Gain Score'])
gains_kb = pd.concat([cols,imp],axis=1)
gains_kb
```

```
[30]:
```

	Feature_Names	Gain Score
0	battery_power	31.190873
1	blue	0.459678
2	clock_speed	0.613019
3	dual_sim	0.489493
4	fc	0.830396
5	four_g	1.167108
6	int_memory	2.965758
7	m_dep	1.644102
8	mobile_wt	3.625213
9	n_cores	2.582847
10	pc	0.874296
11	px_height	19.547371
12	px_width	22.520097
13	ram	3526.232362
14	sc_h	2.203504
15	sc_w	1.597310
16	talk_time	1.666576
17	three_g	0.439621
18	touch_screen	1.458911
19	wifi	0.261677

```
[31]: gains_kb.nlargest(10,'Gain Score')
```

```
[31]:
```

	Feature_Names	Gain Score
13	ram	3526.232362
0	battery_power	31.190873
12	px_width	22.520097
11	px_height	19.547371
8	mobile_wt	3.625213
6	int_memory	2.965758
9	n_cores	2.582847
14	sc_h	2.203504

```
16      talk_time      1.666576
7       m_dep          1.644102
```


```
[32]: gains_kb.nsmallest(10, 'Gain Score')
```

```
[32]:
```

	Feature_Names	Gain Score
19	wifi	0.261677
17	three_g	0.439621
1	blue	0.459678
3	dual_sim	0.489493
2	clock_speed	0.613019
4	fc	0.830396
10	pc	0.874296
5	four_g	1.167108
18	touch_screen	1.458911
15	sc_w	1.597310

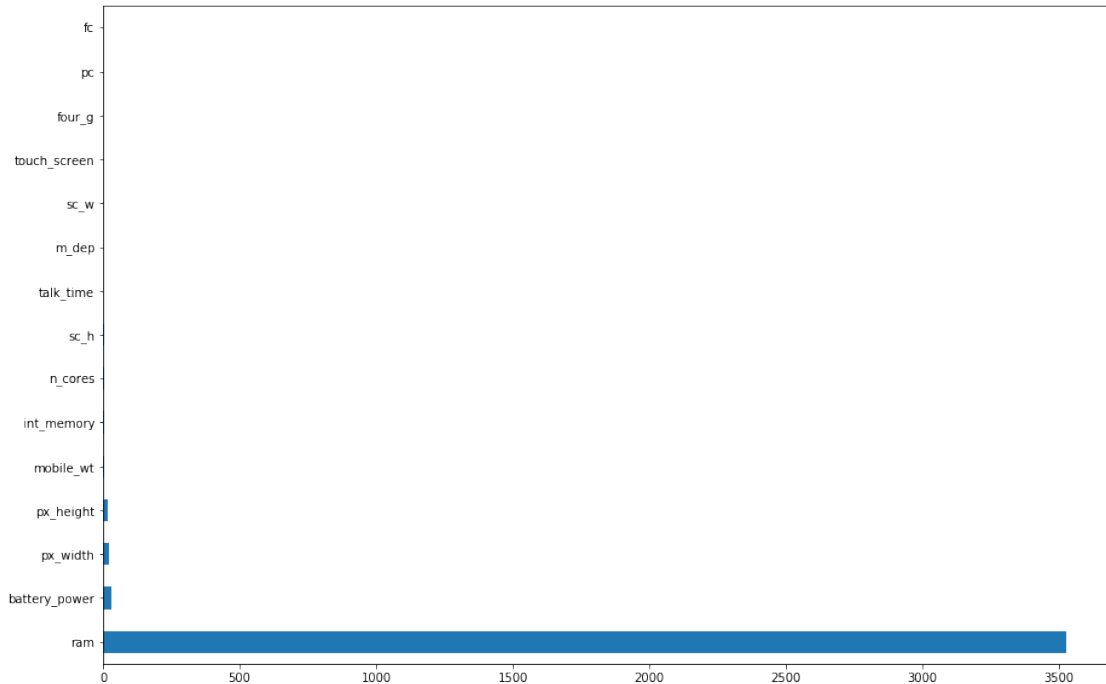
```
[33]: gains_kb.plot(kind='bar',figsize=(10,10))
```

```
[33]: <matplotlib.axes._subplots.AxesSubplot at 0x284695b1d48>
```



```
[34]: features = pd.Series(skb.scores_, index = x.columns)
      features.nlargest(15).plot(kind='barh',figsize=(15,10))
      #plt.savefig('score1.png')
```

```
[34]: <matplotlib.axes._subplots.AxesSubplot at 0x28468c7d8c8>
```



0.2 Principal Component Analysis (PCA)

```
[35]: # Scaling Dataframe
      from sklearn.preprocessing import MinMaxScaler
      mms = MinMaxScaler()
```

```
[38]: x_scaled = mms.fit_transform(x)
```

```
[39]: x_scaled.shape
```

```
[39]: (2000, 20)
```

```
[40]: x.shape
```

```
[40]: (2000, 20)
```

```
[41]: # Applying PCA
      from sklearn.decomposition import PCA
      pca = PCA(n_components=3)
```

```
[42]: x_pca = pca.fit_transform(x_scaled)
```

```
[43]: x_pca
```

```
[43]: array([[ 0.20723861,  0.09565505, -0.00565893],
           [ 0.28947651, -0.00206982,  0.80169178],
           [ 0.93249226, -0.12077773, -0.69947386],
           ...,
           [-0.56557579,  0.03097499,  0.06335067],
           [-0.56346383,  0.76939573, -0.43350877],
           [-0.51957598, -0.66046955, -0.37689763]])
```

```
[44]: features = pd.DataFrame(x_pca, columns=['pca0', 'pca1', 'pca2'])
      features
```

```
[44]:
```

	pca0	pca1	pca2
0	0.207239	0.095655	-0.005659
1	0.289477	-0.002070	0.801692
2	0.932492	-0.120778	-0.699474
3	0.298751	-0.006629	0.774752
4	-0.540990	0.169750	0.815691
...
1995	-0.519601	-0.543636	0.462132
1996	0.293092	-0.554112	-0.381401
1997	-0.565576	0.030975	0.063351
1998	-0.563464	0.769396	-0.433509
1999	-0.519576	-0.660470	-0.376898

[2000 rows x 3 columns]

```
[45]: pip install plotly
```

Requirement already satisfied: plotly in d:\anaconda3\lib\site-packages (5.9.0)
Requirement already satisfied: tenacity>=6.2.0 in d:\anaconda3\lib\site-packages
(from plotly) (8.0.1)
Note: you may need to restart the kernel to use updated packages.

```
[46]: import plotly.express as px
      px.scatter_3d(features, x='pca0', y='pca1', z='pca2')
```

1 KNN

1.1 KNN with ExtraTreeClassifier

```
[82]: from sklearn.neighbors import KNeighborsClassifier
      knn = KNeighborsClassifier()
```

```
[83]: # list of largest scored features
      features_x = list(gains_x.nlargest(15, 'Gain_Score').Feature_Names)
```



```
[84]: x_x = df[features_x]
      x_x.shape
```

```
[84]: (2000, 15)
```

```
[85]: xtrain_x,xtest_x,ytrain_x,ytest_x = train_test_split(x_x,y,train_size=.
      ↪7,random_state=43)
```

```
[86]: knn.fit(xtrain_x,ytrain_x)
```

```
[86]: KNeighborsClassifier()
```

```
[87]: # Testing accuracy
      accuracy_x = knn.score(xtest_x,ytest)
      accuracy_x
```

```
[87]: 0.9166666666666666
```

1.2 KNN with Select k-Best

```
[69]: # list of largest scored features
      features_kb = list(gains_kb.nlargest(15,'Gain Score').Feature_Names)
      features_kb
```

```
[69]: ['ram',
      'battery_power',
      'px_width',
      'px_height',
      'mobile_wt',
      'int_memory',
      'n_cores',
      'sc_h',
      'talk_time',
      'm_dep',
      'sc_w',
      'touch_screen',
      'four_g',
      'pc',
      'fc']
```

```
[71]: x_kb = df[features_kb]
      x_kb.shape
```

```
[71]: (2000, 15)
```

```
[72]: xtrain_kb,xtest_kb,ytrain_kb,ytest_kb = train_test_split(x_kb,y,train_size=.
      ↪7,random_state=43)
```

```
[77]: xtrain_kb.shape
```

```
[77]: (1400, 15)
```

```
[78]: xtest_kb.shape
```

```
[78]: (600, 15)
```

```
[75]: knn = KNeighborsClassifier()  
      knn.fit(xtrain_kb,ytrain_kb)
```

```
[75]: KNeighborsClassifier()
```

```
[80]: # Testing accuracy  
      accuracy_kb = knn.score(xtest_kb,ytest_kb)  
      accuracy_kb
```

```
[80]: 0.9166666666666666
```

1.3 KNN with Principle Component Analysis (PCA)

```
[89]: xtrain_pca, xtest_pca, ytrain_pca, ytest_pca = train_test_split(features_pca,   
      ↪y, test_size=.3, random_state=43)
```

```
[91]: xtrain_pca.shape
```

```
[91]: (1400, 3)
```

```
[92]: xtest.shape
```

```
[92]: (600, 20)
```

```
[90]: knn = KNeighborsClassifier()  
      knn.fit(xtrain_pca,ytrain_pca)
```

```
[90]: KNeighborsClassifier()
```

```
[94]: # Testing accuracy  
      accuracy_pca = knn.score(xtest_pca,ytest_pca)  
      accuracy_pca
```

```
[94]: 0.3
```

2 Comparison

```
[99]: df_comparison = pd.DataFrame([['KNN With_␣  
    ↳ExtraTreesClassifier',accuracy_x],['KNN with Select_␣  
    ↳K-best',accuracy_kb],['KNN with_␣  
    ↳PCA',accuracy_pca]],columns=['Method','Accuracy'])  
df_comparison
```

```
[99]:
```

	Method	Accuracy
0	KNN With ExtraTreesClassifier	0.916667
1	KNN with Select K-best	0.916667
2	KNN with PCA	0.300000

```
[101]: sns.barplot(x='Accuracy',y='Method',data=df_comparison)
```

```
[101]: <matplotlib.axes._subplots.AxesSubplot at 0x2846becbec8>
```

